

Operating Systems

Programming Assignment #5


Working set optimization

Introduction

- Working set refers to the set of pages that a process accesses
- Large working set may result in memory contention and decrease the degree of multiprogramming
- The design of data structure is closely related to the working set size
 - Reducing working set size → improving memory locality

Original Materials


- data.txt
 - Key, value pairs
- test.txt
 - queries
- search.c
 - A program execute queries



max rss = 25476

data.txt

- A large file of sorted records



test.txt format

- <key> <start> <end>\n
- <key> <start > <end>\n
- <key> <start > <end>\n
-

- E.g., f2f5 645 660
- Find key f2f5 and print characters between byte positions <start> and <end>

| index | start | end |
|-------|-------|------|
| 0000 | 2296 | 2490 |
| 978c | 1000 | 4022 |
| a304 | 2613 | 3739 |
| 00ac | 3424 | 3594 |
| 7f65 | 21 | 3422 |
| 1143 | 3732 | 3792 |
| 014d | 2457 | 2537 |
| 7477 | 3382 | 3420 |
| 01b6 | 1021 | 3472 |
| fa4f | 1111 | 3518 |
| 021a | 3653 | 3949 |
| 025c | 1955 | 4049 |
| f2f5 | 745 | 1307 |
| 02df | 3854 | 3865 |
| 0312 | 425 | 2367 |
| 0344 | 2252 | 2421 |
| 0398 | 2556 | 3552 |
| 8ed5 | 2174 | 2361 |
| 03e1 | 1385 | 2306 |
| d0ff | 3763 | 3863 |
| 0439 | 3277 | 3912 |
| 8a3b | 1183 | 2364 |
| 048b | 1923 | 2757 |
| 65ad | 2160 | 3052 |
| aab7 | 2318 | 2926 |

search.c

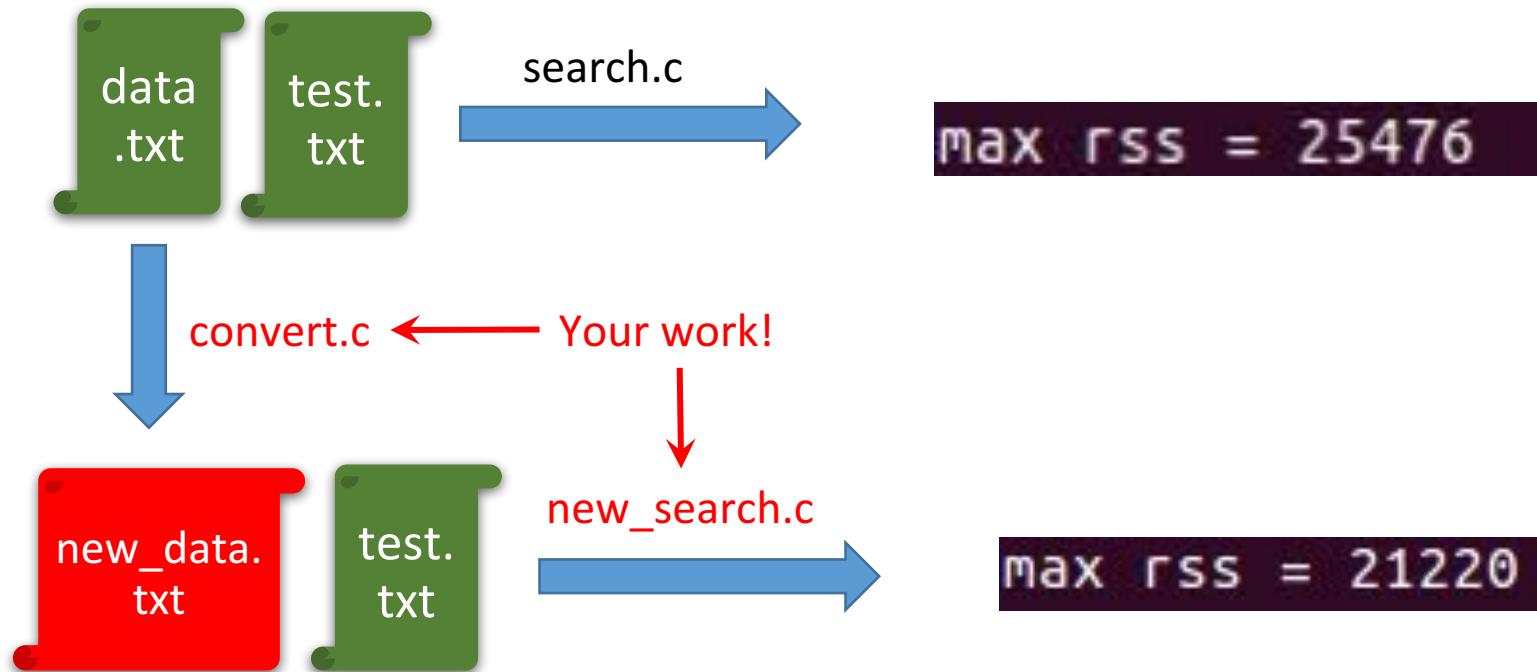
- Map data.txt to memory using mmap()
- Answer queries in test.txt using binary search
- Output:
 - For each query
 - Found → printf("key %s found : %s\n", key, buffer);
 - Not found → printf("key %s not found\n", key);
 - Last output
 - printf("max rss = %lu KB\n", ru_maxrss);

```
key ce33 not found
key fe0f found : 1001100000101101
key fe3e found : 0010100110111111011
key fe7c found : 11101100111101000010
key fec0 found : 1100000000000000
key fef4 found : 10100101001110101
key ff23 found : 001100101010
key ff40 found : 010000100100111100
key ff7d found : 100010110011
key 177c not found
max rss = 46660
```

Problem Statement

- Write **two** programs
- `convert.c`
 - Convert `data.txt` to `new_data.txt`
- `new_search.c`
 - Search `new_data.txt` using the test cases in `test.txt`
- Your new program (`new_search.c`) should have smaller working set compared to the `search.c`!
 - At least 2000 KB reduction

Procedure



RSS (Resident Set Size):

How much physical memory is occupied by a process

convert.c

- Generate new_data.txt
 - Convert data.txt into your own file format
 - Do not modify the keys and values in data.txt!

new_search.c

- Implement your search for new_data.txt
 - Search in the memory space mapped to new_data.txt as the original *search.c* does
 - `mmap()`, `munmap()`
- Output format must be the same as *search.c*
 - Refer to slide No. 6
- Must produce a reduced rss

```
key ce33 not found
key fe0f found : 1001100000101101
key fe3e found : 0010100110111111011
key fe7c found : 11101100111101000010
key fec0 found : 110000000000000
key fef4 found : 10100101001110101
key ff23 found : 001100101010
key ff40 found : 010000100100111100
key ff7d found : 100010110011
key 177c not found
max rss = 46660
```

Useful APIs

- `<resources.h>`
 - `getrusage()`
 - `long ru_maxrss; /* maximum resident set size , unit : KB*/`
- `<sys/mman.h>`
 - `mmap()`, `munmap()`
 - Mapping a file into the logical memory space

Grading Policy

- Output Correctness
 - Must be the same format as search.c (slide No. 6)
 - TAs will test your code with another testcase.
- Search in the memory space mapped to new_data.txt
 - mmap(), munmap()
- RSS reduction
 - ru_max rss should be decreased by at least 2000KB!

Grading Policy

- Uploaded File to E3
 - `student_number.zip`
 - `convert.c/cpp`
 - `new_search.c/cpp`
 - Please follow the naming convention
 - Wrong file formats incur -10 pts score penalty
- Do not plagiarize

Testing OS Environment

- Ubuntu 16.04, Ubuntu 14.04 or CS linux work station
 - `gcc convert.c -o convert`
 - `gcc new_search.c -o new_search`
- Your code should compile successfully in one of the above OS environments